

**DataStax**

# Building resilient and scalable API backends with Apache Pulsar and Spring Reactive

Presentation at ApacheCon@Home 2021  
September 21, 2021

**Lari Hotari** @lhotari  
Senior Software Engineer, DataStax, Inc

Apache Pulsar committer

# Who



Lari Hotari @lhotari

Software Engineer @DataStax, working on Luna Streaming powered by Apache Pulsar

Open-source contributor, Apache Pulsar committer

# Agenda

**01**

Pulsar & Reactive  
Messaging

**02**

Scaling performance  
of message  
processing

**03**

Sample use case &  
Live coding

# Agenda

**01**

Pulsar & Reactive  
Messaging

**02**

Scaling performance  
of message  
processing

**03**

Sample use case &  
Live coding

# Apache Pulsar – Favourite properties

**01**

Streaming, pub-sub  
and queuing come  
together

**03**

Cloud Native  
Architecture

Layered storage  
architecture

**02**

Many options for  
scaling - not just about  
partitions

First class support  
for Kubernetes

Stateless Brokers

# Reactive Message handling – why?

Fault tolerance with  
**fallbacks, retries and  
timeouts**

Compatibility with  
Reactive Streams to  
achieve **reactive from  
end-to-end**

Performance by  
**parallelism and  
pipelining in  
processing**

Resilience with flow  
control (**backpressure**)

Simplification by  
**data oriented** &  
functional approach

Efficiency by optimal  
resource usage

# Reactive Streams adapter library for Apache Pulsar

- <https://github.com/lhotari/reactive-pulsar> , License: ASL 2.0
- Wraps the Apache Pulsar Java Client async API with a simple and intuitive reactive API
- Goes beyond a wrapper
  - Back-pressure support for provided interfaces
  - Pulsar client resource lifecycle management
  - Message producer caching
  - In-order parallel processing at the application level
  - Spring Boot starter for auto configuring a Pulsar Java client and a ReactivePulsarClient

# Reactive Messaging Application building blocks

provided by the Reactive Pulsar Adapter library

Message Sender

Message Reader

Message Consumer

Message Listener Container

- **Sender** - for sending messages with a specific configuration to a specific destination topic.
- **Reader** - The application decides the position where to start reading messages. Position can be earliest, latest, or an absolute or time based position. Suitable also for short-time message polling.
- **Consumer** - for guaranteed message delivery and handling. The broker redelivers unacknowledged messages. Used for both “always on” use cases and batch processing or short-time message consuming or polling use cases.
- **Message Listener Container** - integrates a consumer with the Spring application lifecycle.

\*) The abstractions of the Reactive Pulsar Adapter library are slightly different from the abstractions in Pulsar Java Client. The sender is one example of a difference. The lifecycles of the abstractions are completely different. The `ReactiveMessageSender`, `ReactiveMessageReader` and `ReactiveMessageConsumer` instances themselves are stateless. “Nothing happens until you subscribe” - the key difference in the Reactive programming model.



# Basics of Reactive APIs

- “Nothing happens until you subscribe”
  - Assembly-time vs Runtime
- API calls return a reactive publisher type
  - `Flux<T>` or `Mono<T>` when using Project Reactor
  - `Mono<Void>` for calls that don’t return data

# Reactive Pulsar Adapter

```
public interface ReactivePulsarClient {  
  
    <T> ReactiveMessageSenderBuilder<T> messageSender(Schema<T> schema);  
  
    <T> ReactiveMessageReaderBuilder<T> messageReader(Schema<T> schema);  
  
    <T> ReactiveMessageConsumerBuilder<T> messageConsumer(Schema<T> schema);  
  
}
```

```
public interface ReactiveMessageSender<T> {  
  
    Mono<MessageId> sendMessage(Mono<MessageSpec<T>> messageSpec);  
  
    Flux<MessageId> sendMessages(Flux<MessageSpec<T>> messageSpecs);  
  
}
```

```
public interface ReactiveMessageReader<T> {  
  
    Mono<Message<T>> readMessage();  
  
    Flux<Message<T>> readMessages();  
  
}
```

```
public interface ReactiveMessageConsumer<T> {  
  
    <R> Mono<R> consumeMessage(  
        Function<Mono<Message<T>>, Mono<MessageResult<R>>> messageHandler);  
  
    <R> Flux<R> consumeMessages(  
        Function<Flux<Message<T>>, Flux<MessageResult<R>>> messageHandler);  
  
}
```

# Agenda

**01**

Pulsar & Reactive  
Messaging

**02**

Scaling performance  
of message  
processing


**03**

Sample use case &  
Live coding

**All significant improvements in  
system scalability come from  
parallelism.**

# Pipelining

HTTP/ 1.1 protocol pipelining illustration



“This idea is an extension of the idea of parallelism. It is essentially handling the activities involved in instruction execution as an assembly line. As soon as the first activity of an instruction is done you move it to the second activity and start the first activity of a new instruction. This results in executing more instructions per unit time compared to waiting for all activities of the first instruction to complete before starting the second instruction.”

<https://www.d.umn.edu/~gshute/arch/great-ideas.html>

# In-order parallel processing strategies

- System level: Multiple topic partitions
  - Message is mapped to a specific partition by hashing the key

```
partition = hash(message key) % number_of_partitions
```
- Application level: Message consumer parallelism
  - Micro-batching
    - Messages are consumed in batches which is limited by time and number of entries.
    - The processing can sort and group the entries for parallel processing.
  - Stream splitting / routing / grouping
    - Message is mapped to a specific processing group based on the hash of the message key

```
processing_group = hash(message key) % number_of_processing_groups
```
    - This is well suited for splitting the stream with Project Reactor's `.groupBy`
    - Benefit over micro-batching: no extra latency caused by batching

# Parallel processing with ReactiveMessageConsumer

- An enabler is the special signature of the consumeMessages method on the ReactiveMessageConsumer interface.
  - The method accepts a function that takes `Flux<Message<T>>` input and produces `Flux<MessageResult<R>>>` output.
  - MessageResult combines acknowledgement and the result. The implementation consumes the `Flux<MessageResult<R>>>`, handles the acknowledgements, and returns a `Flux<R>`.
  - This enables guaranteed message delivery combined with stream transformations.

```
public interface ReactiveMessageConsumer<T> {  
    <R> Flux<R> consumeMessages(  
        Function<Flux<Message<T>>, Flux<MessageResult<R>>>>  
        messageHandler);  
}  
  
public interface MessageResult<T> {  
    static <T> MessageResult<T> acknowledge(MessageId messageId, T value) {}  
    static <T> MessageResult<T> negativeAcknowledge(MessageId messageId, T  
    value) {}  
    static MessageResult<Void> acknowledge(MessageId messageId) {}  
    static MessageResult<Void> negativeAcknowledge(MessageId messageId) {}  
    static <V> MessageResult<Message<V>> acknowledgeAndReturn(Message<V>  
    message) {}  
    boolean isAcknowledgeMessage();  
    MessageId getMessageId();  
    T getValue();  
}
```

# Agenda

**01**

Pulsar & Reactive  
Messaging

**02**

Scaling performance  
of message  
processing

**03**

Sample use case &  
Live coding



# Sample use cases – simplified IoT backend

<https://github.com/lhotari/reactive-pulsar-showcase> for complete sample (work in progress)

Telemetry ingestion  
from IoT devices

Telemetry processing

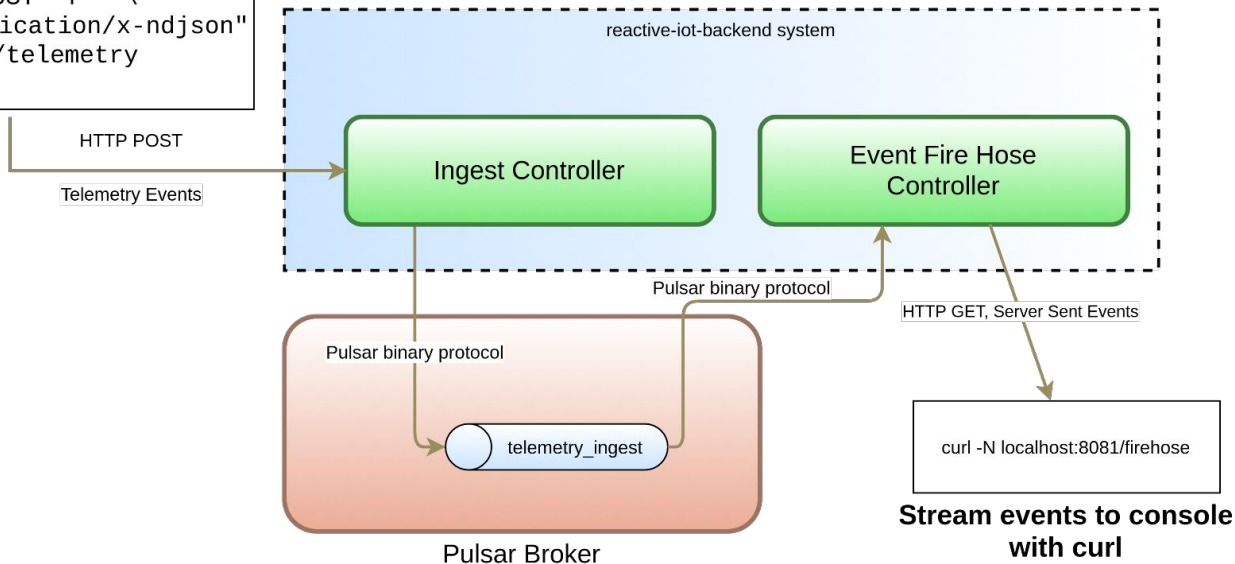
Sending alerts to  
external application - webhook  
interface

Streaming of telemetry values  
to other applications

# Our goal for live coding

**Generate 1 million telemetry events  
with a shell script and curl**

```
{ for i in {1..1000000}; do echo '{"n":  
"device'$i'/sensor1", "v": '$i'.123}'; done;  
} \  
| curl -X POST -T - \  
-H "Content-Type: application/x-ndjson"  
localhost:8081/telemetry
```



Source code: <https://github.com/lhotari/reactive-iot-backend-ApacheCon2021>

```

@RestController
public class IngestController {
    private final ReactiveMessageSender<TelemetryEvent> reactiveMessageSender;

    public IngestController(ReactivePulsarClient reactivePulsarClient,
                           ReactiveProducerCache reactiveProducerCache) {
        reactiveMessageSender = reactivePulsarClient
            .messageSender (Schema.JSON(TelemetryEvent.class))
            .topic ("telemetry_ingest")
            .cache(reactiveProducerCache)
            .maxInflight (100)
            .build();
    }

    @PostMapping("/telemetry")
    public Mono<Void> ingestTelemetry(@RequestBody Flux<TelemetryEvent> telemetryEventFlux) {
        return reactiveMessageSender
            .sendMessages (telemetryEventFlux.map(MessageSpec::of))
            .then();
    }
}

```

```

@RestController
public class EventFireHoseController {
    private final ReactiveMessageReader<TelemetryEvent> reactiveMessageReader;

    public EventFireHoseController(ReactivePulsarClient reactivePulsarClient) {
        reactiveMessageReader = reactivePulsarClient
            .messageReader (Schema.JSON(TelemetryEvent.class))
            .topic ("telemetry_ingest")
            .startAtSpec (StartAtSpec.ofEarliest())
            .endOfStreamAction (EndOfStreamAction.POLL)
            .build();
    }

    @GetMapping("/firehose")
    public Flux<ServerSentEvent<TelemetryEvent>> firehose() {
        return reactiveMessageReader
            .readMessages()
            .map(message -> ServerSentEvent.builder(message.getValue())
                .id(message.getMessageId().toString())
                .build());
    }
}

```

# References

**Apache Pulsar:** <https://pulsar.apache.org/>

**Spring Reactive:** <https://spring.io/reactive>

**Reactive Pulsar adapter:** <https://github.com/lhotari/reactive-pulsar>

Status: *experimental* version 0.1.0 available for use, API is subject to change until 1.0 is released.

**Reactive Pulsar showcase application:** <https://github.com/lhotari/reactive-pulsar-showcase>

Status: work in progress, demonstrates the usage of the Reactive Pulsar Adapter.

**Live coding source code:** <https://github.com/lhotari/reactive-iot-backend-ApacheCon2021>

For usage questions, please use [apache-pulsar](#) and [reactive-pulsar](#) tags on Stackoverflow.

DS

# Thank you !

We are hiring: <https://www.datastax.com/company/careers>